



THE GEORGE  
WASHINGTON  
UNIVERSITY  
WASHINGTON, DC

Introduction to Big Data & Analytics

Instructor: Professor Abzetskhan Adamov

Group 7

Students: Rustam Talibzade, Emin Alizada

1. This project is about learning and practicing Graph Analytics. We were given a huge dataset of files with messages of people in Enron Corporation. As the data is too large, we were allowed to subset this data. Our group was assigned to subset the data by using all top-folders but limiting sub-folders with first 8 and then retrieve all mails within them. To accomplish parsing of given e-mails, we have used two R packages – **readr**, **stringr**. Readr is used for reading the context of the given files, and from readr we have used read\_file function which reads all the contents of file. Stringr is used for string manipulations, to be more precise we have used following functions from stringr:

str\_locate – to find position of given substring in string

str\_sub – to get substring of string

str\_replace\_all – to replace given substring with other one.

After reading all top folders and first 8 subfolders of each top folder we have received approximately 141000 row of records.

After doing some cleanup such as removing rows having NA values and duplicated rows, we have 80860 records.

df1 80860 obs. of 3 variables

2. To reach the goals of our project we installed igraph package:

```
> install.packages("igraph")
WARNING: Rtools is required to build R packages but is not currently installed. Please d
ownload and install the appropriate version of Rtools before proceeding:

https://cran.rstudio.com/bin/windows/Rtools/
Installing package into 'C:/Users/user/Documents/R/win-library/4.1'
(as 'lib' is unspecified)
trying URL 'https://cran.rstudio.com/bin/windows/contrib/4.1/igraph_1.2.11.zip'
Content type 'application/zip' length 9014721 bytes (8.6 MB)
downloaded 8.6 MB


package 'igraph' successfully unpacked and MD5 sums checked

The downloaded binary packages are in
  C:\Users\Public\Documents\wondershare\CreatorTemp\RtmpGk2IB9\downloaded_packages
```

3. After installing igraph package we turned our data frame in igraph object and tried to plot it. As expected, we got a very messy blue bulb.




If we remove labels for simpler visuality:



Because of the fact that dataset is huge even after doing some cleanup, in order to work on this dataset on the local machine we needed to simplify it further. Firstly, we decided to remove all the mails outside of corporation, which means mail addresses that don't end with enron.com. Later we removed mails where from and to is the same (which means mails sent to yourself). The most important simplification was grouping the data by fields from, to and assigning weight according how frequent these email addresses communicate. As the final step we also used igraph's built in simplify function. As a result, following plots were produced.



If we remove labels for simpler visuality:



\* We have used string representation of graph with `str()` function.

\* We have used `igraph::get.adjacency` function like: **`igs.adj <- igraph::get.adjacency(igs)`** and printed first 9 rows and columns. Here is small part of output.

```
adnan.patel@enron.com      a..shankman@enron.com actforchange.com@mailman.enron.com adam.johnson@enron.com
40ees@enron.com            .                               .                               .
a..allen@enron.com          .                               .                               .
a..lindholm@enron.com       .                               .                               .
a..martin@enron.com         .                               .                               .
a..shankman@enron.com       .                               .                               .
actforchange.com@mailman.enron.com .                               .                               .
adam.johnson@enron.com      1                               .                               .
administration.enron@enron.com .                               .                               .
adnan.patel@enron.com      .                               .                               .
```

\* To get list of vertexes we have used `V` function like: `V(igs)`

\* To get list of edges we have used `E` function like: `E(igs)`. Small portion of output (arrows show direction)

```

E(igs)
28639/28639 edges from 170209e (vertex names):
[1] 40ees@enron.com      ->bob.deitz@enron.com      40ees@enron.com      ->steve.wurzel@enron.com
[3] a..allen@enron.com   ->katina.smith@enron.com   a..lindholm@enron.com->john.lamb@enron.com
[5] a..lindholm@enron.com->m..presto@enron.com   a..lindholm@enron.com->mike.curry@enron.com
[7] a..lindholm@enron.com->michael.payne@enron.com a..martin@enron.com ->david.baumbach@enron.com
[9] a..martin@enron.com ->edward.gottlob@enron.com a..martin@enron.com ->elsa.villarreal@enron.co
[11] a..martin@enron.com ->eric.bass@enron.com a..martin@enron.com ->greg.mcclendon@enron.com
[13] a..martin@enron.com ->j..farmer@enron.com a..martin@enron.com ->jim.schwieger@enron.com
[15] a..martin@enron.com ->metz.carey@enron.com a..shankman@enron.com->david.oxley@enron.com
[17] a..shankman@enron.com->greg.whalley@enron.com a..shankman@enron.com->rick.buy@enron.com
[19] a..shankman@enron.com->s..bradford@enron.com a..shankman@enron.com->ben.glisan@enron.com
... omitted several edges

```

\* To get density of graph we used function from SNA library `igs.density = sna::gden(igs.adj)` which accepts adjacency matrix.

\*To get edge density we have used **`igraph::edge_density(igs)`** (which was very small)

```

> igraph::edge_density(igs)
[1] 0.0005361686

```

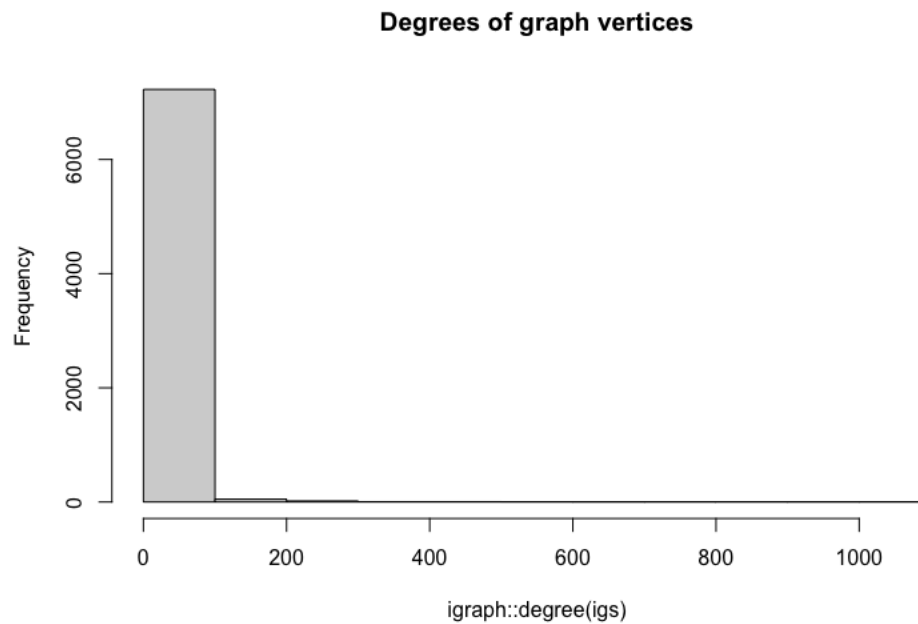
\* To get the degree of vertices we have used **`igraph::degree(igs)`**. Small potion of output

```

              7              5              63
richard.sanders@enron.com      richard.shapiro@enron.com      richard.tomaski@enron.com
              407              75              5
rick.bergsieker@enron.com      rick.buy@enron.com      rick.cates@enron.com
              4              68              2
rika.imai@enron.com      rishi.mehta@enron.com      rita.bahner@enron.com
              52              1              6
rita.wynne@enron.com      rob.bradley@enron.com      rob.brown@enron.com
              16              10              5
rob.cole@enron.com      rob.walls@enron.com      robert.badeer@enron.com
              10              24              64
robert.eickenroht@enron.com      robert.frank@enron.com      robert.gerry@enron.com
              240              14              4
robert.hemstock@enron.com      robert.johnston@enron.com      robert.jones@mailman.enron.com
              22              70              1
robert.knight@enron.com
              6

```

\* We can also visualize this by `hist(igraph::degree(igs), main = "Degrees of graph vertices")`:



\* **To get *betweenness centrality* we have used `igraph::centr_betw(igs)`**

\* We find out degree of a graph with `igraph::diameter(igs)`

```
> igraph::diameter(igs)
[1] 69
```

4. Functions from igraph.pdf:

- a.) **neighbors** function allows us to find neighbours of our vertices – vertices from which edges are directed to the given node (in) or vertices to which edges that go out from the given vertex point to. This is quite useful for our dataset as we can find to whom the given email address mailed (out) or who mailed to the selected email address (in):

```
> neighbors(igs, 2, "in")
+ 22/7309 vertices, named, from 8064962:
[1] andrea.dahlke@enron.com      carol.lapsley@enron.com      david.forster@enron.com      david.oxley@enron.com
[5] debra.bailey@enron.com       duong.luu@enron.com          heather.choate@enron.com     julie.clyatt@enron.com
[9] l..denton@enron.com          laurel.bolt@enron.com        m.hall@enron.com             madhup.kumar@enron.com
[13] mike.croucher@enron.com      r..harrington@enron.com     s..theriot@enron.com        sonia.hennessy@enron.com
[17] steve.nat@enron.com          tammie.schoppe@enron.com    tara.piazze@enron.com       veronica.espinoza@enron.com
[21] w..white@enron.com           william.crooks@enron.com
```

```
> neighbors(igs, 2, "out")
+ 1/7309 vertex, named, from 8064962:
[1] katina.smith@enron.com
```

- b.) **Is\_directed** function returns true or false depending on whether the graph is directed or not. Here we used the function for our directed graph, as we can see the returned values is TRUE:

```
> is_directed(igs)
[1] TRUE
```

- c.) **Is\_igraph** function returns true or false depending on whether the given variable is igraph. While it seems simple, this function is helpful when testing new functions and it also helped us during this project.

From the screenshot below, we got true for our igraph:

```
> is_igraph(igs)
[1] TRUE
```

- d.) **Gsize()** – returns number of edges in our graph:

```
> gsize(ig)
[1] 80860
> gsize(igs)
[1] 28639
```

e.) **Gorder()** – returns number of vertices in our graph:

```
> gorder(igs)
[1] 14322
> gorder(igs)
[1] 7309
>
```

f.) **random\_walk** – this function starts from the mentioned vertices which is a start point and randomly moves to different vertices X times (where X in our case is 6) if possible (f.e. if let's say the 4<sup>th</sup> node does not point to any other node then the function stops):

```
> random_walk(igs, 21, 6)
+ 6/7309 vertices, named, from 693ee03:
[1] alex.perwich@enron.com      greg.whalley@enron.com      john.lavorato@enron.com      greg.whalley@enron.com
[5] christie.patrick@enron.com   jeff.skilling@enron.com
```

g.) **articulation\_points** – this function returns those vertices that if removed will make the graph more connected:

```
> articulation_points(igs)
+ 344/7309 vertices, named, from 693ee03:
[1] clickathome@enron.com      andrew.lewis@enron.com      drew.fossum@enron.com
[4] jeff.skilling@enron.com     center.dl-portland@enron.com  chris.wiebe@enron.com
[7] burton.mcintyre@enron.com   colleen.koenig@enron.com     larry.lawyer@enron.com
[10] dl-ga-all_enron_houston@enron.com chairman.ken@enron.com      administration.enron@enron.com
[13] dl-ga-all_enron_worldwide1@enron.com announcements.enron@enron.com energy.dl-ga-all_ubsw@enron.com
[16] houston.dl-ubsw@enron.com   infrastructure.ubsw@enron.com georgianne.hodges@enron.com
[19] lindon.chiu@enron.com       preston.ochsner@enron.com    james.noles@enron.com
[22] bob.lee@enron.com           shirley.crenshaw@enron.com   william.kendrick@enron.com
[25] marie.hejka@enron.com       ozzie.pagan@enron.com        sam.romero@enron.com
[28] carol.lapsley@enron.com     domingo.drakes@enron.com     vince.kaminski@enron.com
```

h.) **mean\_distance** – this function returns average value of all shortest paths between vertices.  
This is what we got for our graph:

```
> mean_distance(igs)
[1] 4.330467
```

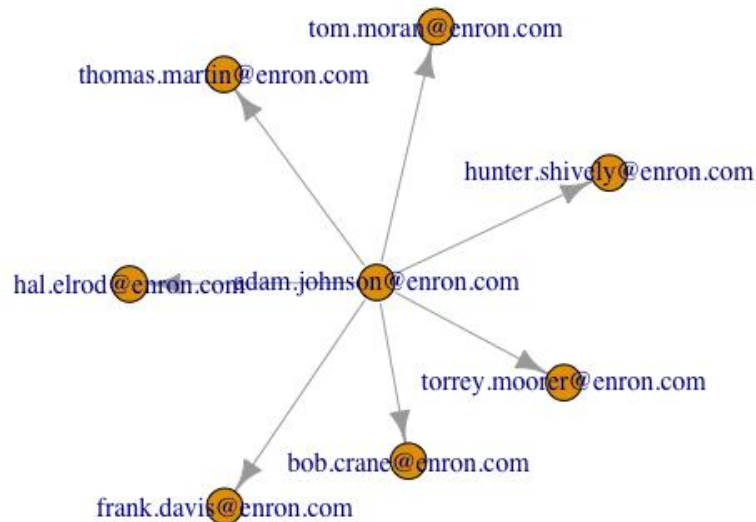


- i.) **igraph::transitivity** - Transitivity measures the probability that the adjacent vertices of a vertex are connected. This is sometimes also called the clustering coefficient.

```
> igraph::transitivity(igs)
[1] 0.03923655
```

- j.) Using **subgraph.edges** function we can get subgraph of our graph with specified edges. In our dataset it can be useful to find the connections for the given email addresses. We plotted our new subgraph using the piece of code mentioned below:

```
plot.igraph(subgraph.edges(igs, 37:43))
```



5. Determine the (a) central nodes(s) in the graph, (b) longest path(s), (c) largest clique(s), (d) ego(s), (e) power centrality, (f) find communities.

a. Central Nodes - `igraph::centr_degree(igs)` (Part of output, because it's too large)

```
[573] 4 94 32 12 17 22 14 15 39 1 28 11 4 4 10 13 33 16 8 35 31 100
[595] 109 3 29 43 2 12 1 43 22 12 47 373 3 21 9 7 63 14 2 3 73 162
[617] 14 3 12 29 25 127 63 29 89 11 51 39 1 6 14 9 5 408 37 18 5 17
[639] 60 19 19 4 21 9 2 142 1013 54 76 20 25 20 28 29 30 32 22 8 7 4
[661] 59 16 2 1 10 11 2 59 12 2 2 17 10 105 6 2 14 1 10 30 10 6
[683] 28 186 4 116 15 14 9 21 3 15 243 13 3 5 18 10 45 2 1 11 17 17
[705] 19 6 24 234 15 24 8 12 3 133 45 41 11 55 1 7 2 39 3 11 19 7
[727] 7 26 4 46 27 35 7 23 1 1 23 55 16 8 247 81 2 99 17 27 10 2
[749] 3 153 26 193 11 48 60 16 22 129 11 36 20 4 4 10 56 4 2 2 10 7
[771] 35 35 49 8 3 5 1 11 11 8 20 3 20 73 8 25 11 76 36 50 24 8
[793] 36 1 8 30 10 62 5 18 65 17 2 353 2 3 50 19 125 4 12 23 24 75
[815] 5 104 1 9 9 9 12 4 21 17 7 4 5 1 6 6 12 27 4 67 17 10
[837] 66 6 15 15 37 35 50 12 13 10 16 59 55 1 3 1 74 15 11 31 1 1
[859] 1 8 25 5 32 56 7 24 23 4 15 12 88 15 19 25 3 29 34 12 2 171
[881] 3 9 3 323 3 7 2 10 1 113 3 12 8 1 12 28 1 10 50 22 27 25
[903] 3 2 59 3 1 2 4 268 5 44 13 98 16 9 12 6 9 10 28 20 2 13
[925] 2 2 28 40 40 1 1 19 5 4 1 31 5 241 100 3 14 9 7 35 32 15
[947] 2 20 15 33 3 6 10 7 9 3 3 2 60 10 31 6 6 10 2 2 2 84
[969] 5 4 12 224 20 8 2 7 5 65 407 75 5 4 68 2 52 1 6 16 10 5
[991] 10 24 64 240 14 4 22 70 1 6
[ reached getOption("max.print") -- omitted 6309 entries ]

$centralization
[1] 0.06878085

$theoretical_max
[1] 106813728
```

b. Longest path

```
igraph::diameter(igs)
```

```
igraph::farthest_vertices(igs)
```

```
> igraph::diameter(igs)
[1] 69
> igraph::farthest_vertices(igs)
$vertices
+ 2/7309 vertices, named, from 693ee03:
[1] diane.cutsforth@enron.com dan.bump@enron.com

$distance
[1] 69
```

c. Largest clique - `igraph::largest_cliques(igs)`

```
> igraph::largest_cliques(igs)
[[1]]
+ 13/7309 vertices, named, from 693ee03:
[1] kenneth.lay@enron.com      vanessa.groscrand@enron.com  rosalee.fleming@enron.com    katherine.brown@enron.com
[5] sherri.sera@enron.com      tori.wells@enron.com         bobbie.power@enron.com      nicki.daw@enron.com
[9] rex.rogers@enron.com       john.sherriff@enron.com      greg.whalley@enron.com      mark.frevert@enron.com
[13] steven.kean@enron.com

[[2]]
+ 13/7309 vertices, named, from 693ee03:
[1] kenneth.lay@enron.com      vanessa.groscrand@enron.com  rosalee.fleming@enron.com    katherine.brown@enron.com
[5] sherri.sera@enron.com      tori.wells@enron.com         bobbie.power@enron.com      nicki.daw@enron.com
[9] rex.rogers@enron.com       david.delainey@enron.com      greg.whalley@enron.com      mark.frevert@enron.com
[13] steven.kean@enron.com
```

d. Ego - `igraph::ego(igs)`

```
> igraph::ego(igs)[1:10]
[[1]]
+ 3/7309 vertices, named, from 693ee03:
[1] 40ees@enron.com      bob.deitz@enron.com      steve.wurzel@enron.com

[[2]]
+ 24/7309 vertices, named, from 693ee03:
[1] a..allen@enron.com      andrea.dahlke@enron.com    carol.lapsley@enron.com      david.forster@enron.com
[5] david.oxley@enron.com    debra.bailey@enron.com     duong.luu@enron.com          heather.choate@enron.com
[9] julie.clyatt@enron.com   katina.smith@enron.com     l..denton@enron.com          laurel.bolt@enron.com
[13] m.hall@enron.com         madhup.kumar@enron.com     mike.croucher@enron.com      r..harrington@enron.com
[17] s..theriot@enron.com     sonia.hennessy@enron.com   steve.nat@enron.com          tammie.schoppe@enron.com
[21] tara.piazze@enron.com    veronica.espinosa@enron.com w..white@enron.com          william.crooks@enron.com

[[3]]
+ 5/7309 vertices, named, from 693ee03:
```

e. Power centrality - `igraph::power centrality(igs, loops = F, exponent = 0.9)`

```
> igraph::power centrality(igs, loops = F, exponent = 0.9)
      40ees@enron.com      a..allen@enron.com      a..lindholm@enron.com
      0.015803099          -0.004667143          -0.395589265
      a..martin@enron.com    a..shankman@enron.com    actforchange.com@mailman.enron.com
      0.063212394          1.062004881          0.007901549
      adam.johnson@enron.com  administration.enron@enron.com    adnan.patel@enron.com
      -0.494336709          0.023704648          -0.072680054
      adriane.schultea@enron.com  adrienne.engler@enron.com    aile@mailman.enron.com
      -0.013965213          3.476447640          0.007901549
      airam.arteaga@enron.com    alan.aronowitz@enron.com    alan.chen@enron.com
      1.135799240          -0.562905593          0.595471496
      alan.comnes@enron.com      alan.engberg@enron.com      albert.meyers@enron.com
      -1.452069738          -0.704964881          0.022914493
      aleck.dadson@enron.com      alex.hidalgo@enron.com      alex.perwich@enron.com
      -2.180893073          -3.294402859          0.105274190
      alex.saldana@enron.com      alex.villarreal@enron.com    alhamd.alkhayat@enron.com
```


f. Communities - `igraph::walktrap.community(igs)`

```
> igraph::walktrap.community(igs)
IGRAPH clustering walktrap, groups: 108, mod: 0.61
+ groups:
$`1`
 [1] "a..lindholm@enron.com"          "alan.comnes@enron.com"
 [3] "albert.meyers@enron.com"        "amy.copeland@enron.com"
 [5] "amy.jon@enron.com"             "andrea.bertone@enron.com"
 [7] "anna.mehrer@enron.com"          "bachelor.conf.@enron.com"
 [9] "bernadette.hawkins@enron.com"   "bert.meyers@enron.com"
[11] "beth.perlman@enron.com"         "beverly.stephens@enron.com"
[13] "bill.williams@enron.com"        "brad.alford@enron.com"
[15] "brett.wiggs@enron.com"          "brian.redmond@enron.com"
[17] "bruno.gaillard@enron.com"       "bryan.garrett@enron.com"
```

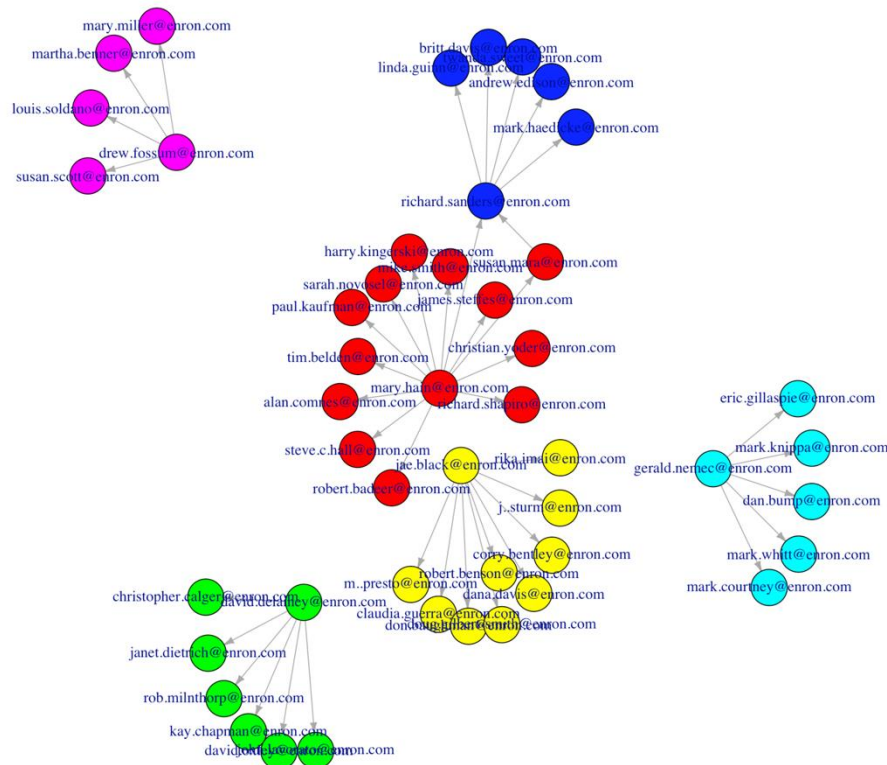
6. Resulting graph with too many vertices and edges will look very messy in the plot. Try to filter vertices and their edges in some way having in resulting plot (visualization) 30 – 100 vertexes. Differentiate vertices (by color, size, shape) and edges (color, type) of graph. Think about opportunity to assign weights to edges differentiating them accordingly.

In order to have around 30-100 vertexes we have used our weights that were assigned earlier, so we decided to filter out all records where weight is less than 33.

```
df2 <- df1[df1$weight > 34, ]
```



The same plot with mail addresses.



Here is the list of final vertexes that we have. Which means they have written to each other the biggest number of emails

```
> V(g2)
+ 47/47 vertices, named, from 58e4533:
[1] david.delainey@enron.com    drew.fossum@enron.com    gerald.nemec@enron.com
[4] jae.black@enron.com         mary.hain@enron.com      richard.sanders@enron.com
[7] susan.mara@enron.com        christopher.calger@enron.com david.oxley@enron.com
[10] janet.dietrich@enron.com    john.lavorato@enron.com  kay.chapman@enron.com
[13] rob.milnthorp@enron.com     louis.soldano@enron.com  martha.benner@enron.com
[16] mary.miller@enron.com       susan.scott@enron.com   dan.bump@enron.com
[19] eric.gillaspie@enron.com    mark.courtney@enron.com mark.knippa@enron.com
[22] mark.whitt@enron.com        claudia.guerra@enron.com corry.bentley@enron.com
[25] dana.davis@enron.com        don.baughman@enron.com  doug.gilbert-smith@enron.com
[28] j..sturm@enron.com          m..presto@enron.com     rika.imai@enron.com
+ ... omitted several vertices
```

Conclusion:

- a. R programming language was relatively new area for both of us, so we started with Tutorial that is uploaded to Blackboard. After getting some essentials of working with the graphs in R with igraph package. We started implementation of the project. Reading the contents of the datasets was the

starting point of this project, so we decided to do it together. We extracted *from*, *to* and *subject* with the help of functions in *stringr* package. Afterwards, we made small research about the ways to simplify data, and answer questions from assignment independently, later combined our ideas, and implemented it altogether. The hardest part of the assignment for us was simplification of data, which took some time for us to do it. We get rid of duplicate data and then used in-built *simplify* function as we mentioned [above](#). For the first time in our lives, we understood how slow our computers can be when they are working with big data. RAM was full, CPU was loaded to 100% and even though it took around 15-20 minutes to construct CSV from the given dataset with the method that was assigned to us. Next, we get acquainted with new functions in *igraph* package and calculated the values we were asked in Question 5. Lastly, we plotted the graph by coloring the vertices to differentiate them.

- b. Adding to what we mentioned above, eventually, we managed to find the email addresses who contacted the most number of people and sent or received the most amount mails among each other.